

097200- Deep learning - HW1

Submission date: 06/12/2020

Theoretical Part:

1. The Softmax function is used to normalize the output of a neural network $f(\cdot; w)$ to a probability distribution over predicted output classes:

$$\hat{y}_i = \text{Softmax}(x)_i = \frac{\exp^{f(x;w)_i}}{\sum_{j=0}^n \exp^{f(x;w)_j}}.$$

Denote $L(\hat{y}, y)$ to be the loss function. Show the derivative of the loss w.r.t the weights. I.e., $\frac{\partial L(\hat{y}, y)}{\partial w}$.

Practical Part:

In the following exercise, you will create a classifier for the MNIST dataset. You should write your own training and evaluation code and meet the following constraints:

- You are only allowed to use torch tensor manipulations.
- You are NOT allowed to use:
 - Auto-differentiation - *backward()*
 - Built-in loss functions
 - Built-in activations
 - Built-in optimization
 - Built-in layers (torch.nn)

The neural network you build should:

- Have at least one hidden layer
- Obtain at least 75% accuracy on the test set

Submission instructions:

Submission **must be in pairs** (course partners) and will contain a short (two pages) pdf report containing:

- Model architecture description, training procedure (hyperparameters, optimization details, etc.).
- Convergence plot of accuracy as a function of time (epochs). The plot should depict both training and test performance (i.e. two curves, one for the train, and one for the test).
- A short summary of your attempts and conclusions.

In addition, you should also supply:

- Code (python file) able to reproduce your results - we might test it on different variants on these datasets.
- The trained network with trained weights (.pkl file).
[the weights tensors can be saved with `torch.save({'w1':w1, 'w2':w2 }, 'path_to_w.pkl')` and load with `torch.load('path_to_w.pkl')`]
- A function called `"evaluate_hw1()"`. The function should load the MNIST test-set, load your trained network (you can assume that the data and model files are located in the script folder), and return the average accuracy over the test-set. This function should be written in a separate script.

Moodle submission:

You should submit a Zip file containing:

- Python files
 - Training procedure
 - Evaluation procedure
- 1 pdf file with
 - Your full names and IDs
 - Typed answers for the theoretical part
 - A short explanation about the practical part
- Pickle file (If the file is too big for the Moodle, upload it to your Google-Drive and copy the link to your pdf report)

Good Luck!